

exLong: Generating Exceptional Behavior Tests with Large Language Models

Jiyang Zhang¹, Yu Liu¹, Pengyu Nie², Junyi Jessy Li¹, and Milos Gligoric¹



ICSE 2025

partially
supported by



Exceptional Behavior Tests (EBTs)

- **Exceptions**: thrown during program execution if an unwanted event happens
- **Exceptional behavior tests**: check that the code detects unwanted events and throws appropriate exceptions

MUT

```
public static int[] qualityScores(final Fastq fastq, final int[] qualityScores) {  
    if (fastq == null) {  
        throw new IllegalArgumentException("fastq must not be null");  
    }  
    if (qualityScores == null) {  
        throw new IllegalArgumentException("qualityScores must not be null");  
    }  
    ...  
}
```

guard (if statement)  **if (fastq == null) {**
exception (throw statement)  **throw new IllegalArgumentException("fastq must not be null");**

EBT

```
@Test  
public void testQualityScoresIntArrayNullFastq() {  
    try {  
        FastqTools.qualityScores(null, new int[0]);  
        Assert.fail(  
            "qualityScores(null, int[]) expected IllegalArgumentException");  
    } catch (IllegalArgumentException e) { }  
}
```

Exceptions Deserve More Test Coverage!

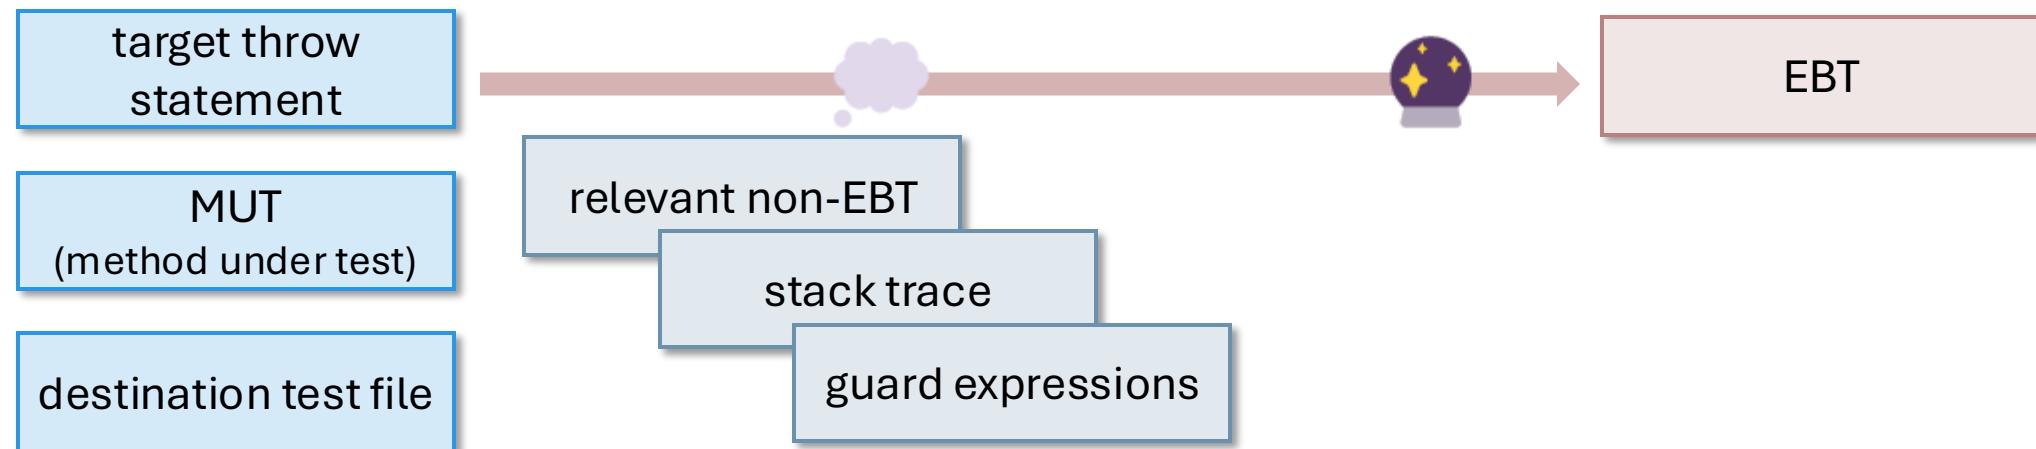
- EBTs check critical error-handling code
- Developers are not writing sufficient EBTs
 - ~40% throw statements covered
 - Developers usually focus on “happy paths” when writing tests [1][2]
- No existing tools specifically help developers write EBTs
- LLMs are struggling to generate correct EBTs

[1] Dalton, Francisco, et al. "Is exceptional behavior testing an exception? an empirical assessment using java automated tests." Proceedings of the 24th International Conference on Evaluation and Assessment in Software Engineering. 2020.

[2] Di Bernardo, Rafael, et al. "Agile testing of exceptional behavior." 2011 25th Brazilian Symposium on Software Engineering. IEEE, 2011.

Our Solution: LLM + Program Analyses

-  exLong
 - base LLM: Code Llama
 - instruction-tuned to reason about context collected via program analyses



Relevant Non-EBT

- Non exception behavior tests that
 - (1) invoke the given MUT, or
 - (2) are in the same destination test file
- Provides example on how to prepare the test inputs & the expected coding style

MUT

```
public static int[] qualityScores(  
    final Fastq fastq,  
    final int[] qualityScores  
) {  
    ...  
}
```

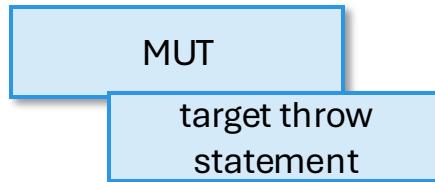
relevant non-EBT

```
@Test  
public void testQualityScoresIntArray() {  
    int[] qualityScores = new int[4];  
    FastqTools.qualityScores(builder.build(), qualityScores);  
    for (int i = 0; i < 4; i++) {  
        Assert.assertTrue(qualityScores[i] != 0);  
    }  
}
```

Stack Trace

- The sequence of method invocations that start from the given MUT and lead to the target throw statement
- Provides insights on what other methods may be involved during execution

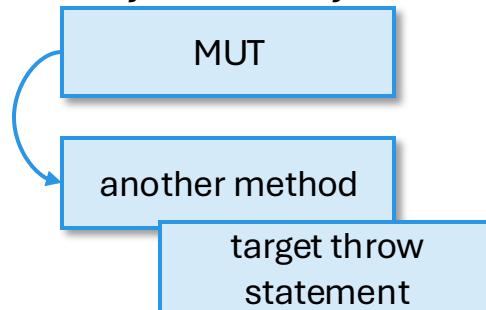
target throw statement is in the MUT



stack trace

qualityScores(FastqTools.java)

*target throw statement is in another method
transitively invoked by the MUT*



searchTerm(SearchCommandParser.java)
parse(SearchCommandParserTest.java)
...

Guard Expressions

- The logical formula representing the constraints on the symbolic variables that must be true to follow the stack trace
- Provides direct guidance on how to trigger the exception

MUT
target throw statement

```
public static int[] qualityScores(final Fastq fastq, final int[] qualityScores) {
    if (fastq == null) {
        throw new IllegalArgumentException("fastq must not be null");
    }
    if (qualityScores == null) {
        throw new IllegalArgumentException("qualityScores must not be null");
    }
    int size = fastq.getQuality().length();
    if (qualityScores.length != size) {
        throw new IllegalArgumentException(
            "qualityScores must be the same length as the FASTQ sequence quality");
    }
    FastqVariant variant = fastq.getVariant();
    for (int i = 0; i < size; i++) {
        char c = fastq.getQuality().charAt(i);
        qualityScores[i] = variant.qualityScore(c);
    }
    return qualityScores;
}
```

guard expressions

```
qualityScores.length !=  
fastq.getQuality().length()
```

Guard Expressions

- The logical formula representing the constraints on the symbolic variables that must be true to follow the stack trace
- Provides direct guidance on how to trigger the exception

MUT

```
public static int[] qualityScores(final Fastq fastq, final int[] qualityScores) {
    if (fastq == null) {
        throw new IllegalArgumentException("fastq must not be null");
    }
    if (qualityScores == null) {
        throw new IllegalArgumentException("qualityScores must not be null");
    }
    int size = fastq.getQuality().length();
    if (qualityScores.length != size) {
        throw new IllegalArgumentException(
            "qualityScores must be the same length as the FASTQ sequence quality");
    }
    FastqVariant variant = fastq.getVariant();
    for (int i = 0; i < size; i++) {
        char c = fastq.getQuality().charAt(i);
        qualityScores[i] = variant.qualityScore(c);
    }
    return qualityScores;
}
```

target throw statement

guard expressions

```
qualityScores.length != fastq.getQuality().length()
```

Use Case: Developer-Oriented

Goal: Generate an EBT for one specific target throw statement



performing program analyses

Generate a test in \${destination test file}
that covers \${target throw statement}
from \${MUT}



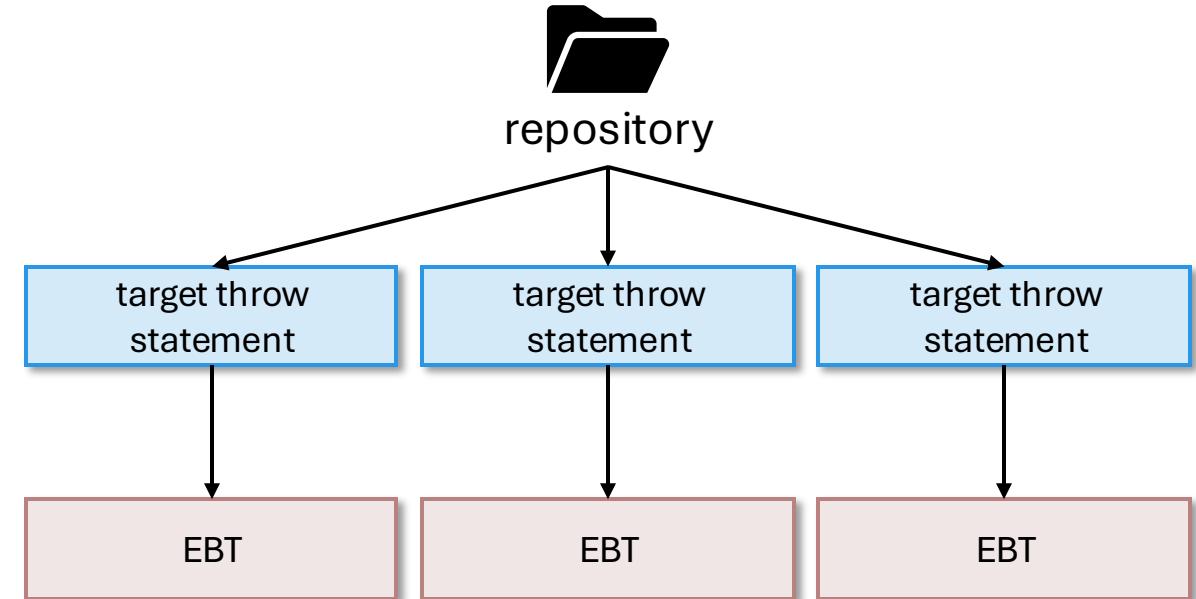
```
@Test  
public void testQualityScoresIntArrayNullFastq() {  
    try {  
        FastqTools.qualityScores(null, new int[0]);  
        Assert.fail(  
            "qualityScores(null, int[]) expected IllegalArgumentException");  
    } catch (IllegalArgumentException e) { }  
}
```

Use Case: Machine-Oriented

Goal: Generate EBTs to cover as many throw statements as possible



- searching for throw statements
- performing program analyses
- generating EBT



Evaluation Dataset

- Open-source Java repositories on GitHub
 - based on CodeSearchNet's list of repositories and cross-project split
 - compile successfully and do not contain test failures
 - have at least one EBT

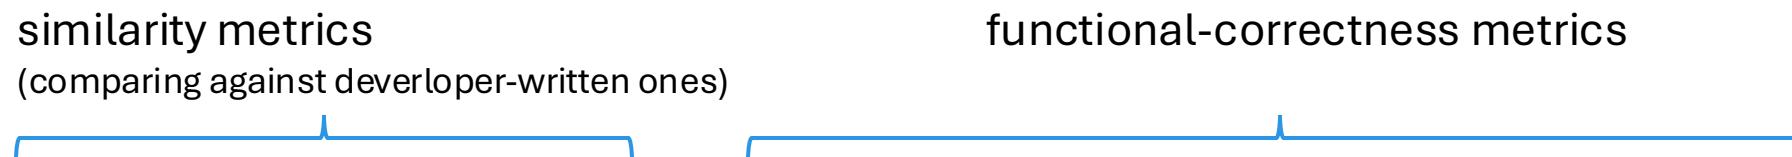
	#Repos	#Tests	#EBTs	
All	562	111,230	12,574	
Train	501	100,030	11,182	
Valid	29	5,298	550	
Eval	32	5,902	842	

instruction-tuning exLong

evaluation on the two use cases

Results: Developer-Oriented

- 434 EBTs, 278 throw statements (covered by at least one non-EBT), 41 exception types



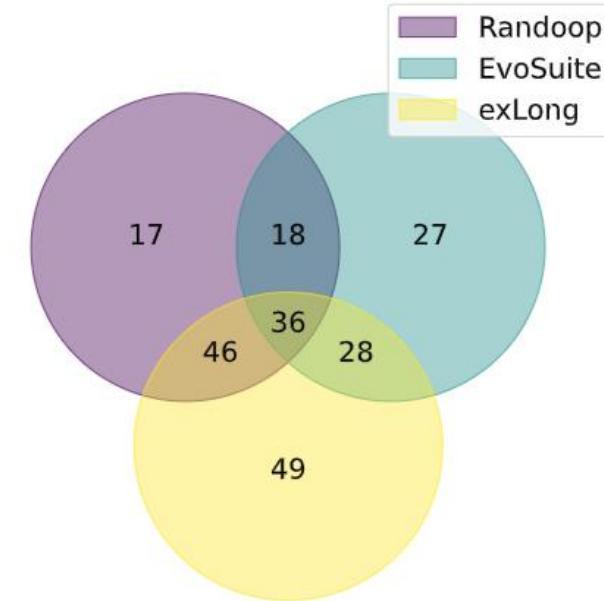
	xMatch	CodeBLEU	Compilable%	Runnable%	ThrowCov%
GPT-3.5	14.98	64.28	75.12	61.29	48.39
CAT-LM	9.83	59.79	71.83	36.64	30.03
exLong w/o context	12.06	60.62	61.52	46.70	38.25
exLong	19.05	67.49	82.10	67.63	59.45

Results: Machine-Oriented

- 649 throw statements, 81 exception types

random / search-based
test generation tools

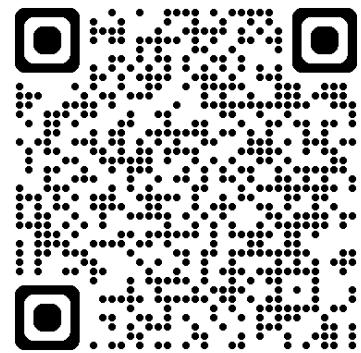
ThrowCov%	
Randoop	21.87
EvoSuite	20.37
exLong	29.72



Conclusions

-  exLong: LLM + program analyses for generating exceptional behavior tests
- Sent 7 PRs with 35 EBTs generated by exLong; 4 accepted, 3 pending
- Tool paper at FSE'25

- Code: <https://github.com/EngineeringSoftware/exlong>
- Data: <https://huggingface.co/datasets/EngineeringSoftware/exLong-dataset>
- Model: <https://huggingface.co/EngineeringSoftware/exLong>



Jiyang Zhang, Yu Liu, **Pengyu Nie**, Junyi Jessy Li, and Milos Gligoric

<jiyang.zhang@utexas.edu> <pynie@uwaterloo.ca>

Backup Slides

Results: exLong based on Proprietary LLM

	xMatch	CodeBLEU	Compilable%	Runnable%	ThrowCov%
GPT-4o	16.82	65.56	81.87	71.52	55.53
exLong-GPT-4o	17.74	66.77	82.49	75.35	64.75

Case Study / Real-World Impact

- Prepare PRs with the EBTs generated by exLong for 9 actively-maintained repositories in the eval set
- For 2 repositories, the same EBTs were added by developer in a later commit (after we collected data)
- Submitted PRs to 7 repositories with 35 EBTs
 - 4 accepted; 1 of them merged within 30min
 - 3 pending

Exceptions

- **Exceptions:** thrown during program execution if an unwanted event happens
 - e.g., invalid input, illegal state
 - “throw” in Java/C++, “raise” in Python

```
public static int[] qualityScores(final Fastq fastq, final int[] qualityScores) {
    if (fastq == null) {
        throw new IllegalArgumentException("fastq must not be null");
    }
    if (qualityScores == null) {
        throw new IllegalArgumentException("qualityScores must not be null");
    }
    int size = fastq.getQuality().length();
    if (qualityScores.length != size) {
        throw new IllegalArgumentException(
            "qualityScores must be the same length as the FASTQ sequence quality");
    }
    FastqVariant variant = fastq.getVariant();
    for (int i = 0; i < size; i++) {
        char c = fastq.getQuality().charAt(i);
        qualityScores[i] = variant.qualityScore(c);
    }
    return qualityScores;
}
```